

Regularized Anderson Acceleration for Off-Policy Deep Reinforcement Learning



Wenjie Shi, Shiji Song, Hui Wu, Ya-Chu Hsu, Cheng Wu, Gao Huang
Tsinghua University, Beijing, China
Email: shiwj16@mails.tsinghua.edu.cn



MOTIVATION

Sample inefficiency


- An enormous number of trials
- Long training time



AlphaGo Zero
millions of self-play



StarCraft
several days



Autonomous Driving
real physical scenarios

Existing methods

- To learn the model of system dynamics.
- To reuse past experience (off-policy).

Main observations

- RL is closely linked to fixed-point problem.
- Anderson acceleration is a method capable of speeding up the computation of fixed point iterations.

METHOD

Fixed-point problem:

Given $g: \mathbb{R}^n \rightarrow \mathbb{R}^n$, solve $x = g(x)$.

RL problem:

Given the optimality Bellman operator \mathcal{T} : solve $Q^\pi = \mathcal{T}Q^\pi$.

Algorithm FPI. Fixed-Point Iteration:

For $k = 0, 1, \dots$
Set $x_{k+1} = g(x_k)$.

Algorithm PI. Policy Iteration:

For $k = 0, 1, \dots$
Set $Q^{\pi_{k+1}} = \mathcal{T}Q^{\pi_k}$.

Algorithm AA. Anderson Acceleration:

For $k = 0, 1, \dots$
Set $F_k = (f_{k-m+1}, \dots, f_k)$, where $f_i = g(x_i) - x_i$.
Solve $\alpha^k = (\alpha_1^k, \dots, \alpha_m^k)^T$:
 $\min_{\alpha} \|F_k \alpha\|_2$ s.t. $\sum \alpha_i = 1$.
Set $x_{k+1} = \sum_{i=1}^m \alpha_i^k g(x_{k-m+i})$.

Algorithm AA for PI.

For $k = 0, 1, \dots$
Set $\Delta_k = (\delta_{k-m+1}, \dots, \delta_k)$, where $\delta_i = \mathcal{T}Q^{\pi_i} - Q^{\pi_i}$.
Solve $\alpha^k = (\alpha_1^k, \dots, \alpha_m^k)^T$:
 $\min_{\alpha} \|\Delta_k \alpha\|_2$ s.t. $\sum \alpha_i = 1$.
Set $Q^{\pi_{k+1}} = \sum_{i=1}^m \alpha_i^k \mathcal{T}Q^{\pi_{k-m+i}}$.

Challenges

- Sweeping entire state-action space is intractable.
- Function approximation errors are unavoidable.
- The solution may suffer from ill-conditioning.

Algorithm RAA. Regularized AA:

$$\min_{\alpha} \|\tilde{\Delta}_k \alpha\|_2^2 + \lambda \|\alpha\|_2^2$$

Proposition 1: Consider two identical PIs P_1 and P_2 with function approximation. P_2 is implemented with RAA, whereas P_1 is only implemented with AA. Let α^k and $\tilde{\alpha}^k$ be the coefficient vectors of P_1 and P_2 respectively.

$$\|\tilde{\alpha}^k\|_2 \leq \sqrt{\frac{\lambda + \|\tilde{\Delta}_k\|_2^2}{m\lambda}}, \quad \|\tilde{\alpha}^k - \alpha^k\|_2 \leq \frac{\|\tilde{\Delta}_k^T \tilde{\Delta}_k - \Delta_k^T \Delta_k\|_2 + \lambda}{\lambda} \|\alpha^k\|_2$$

EXPERIMENTS

Comparative evaluation

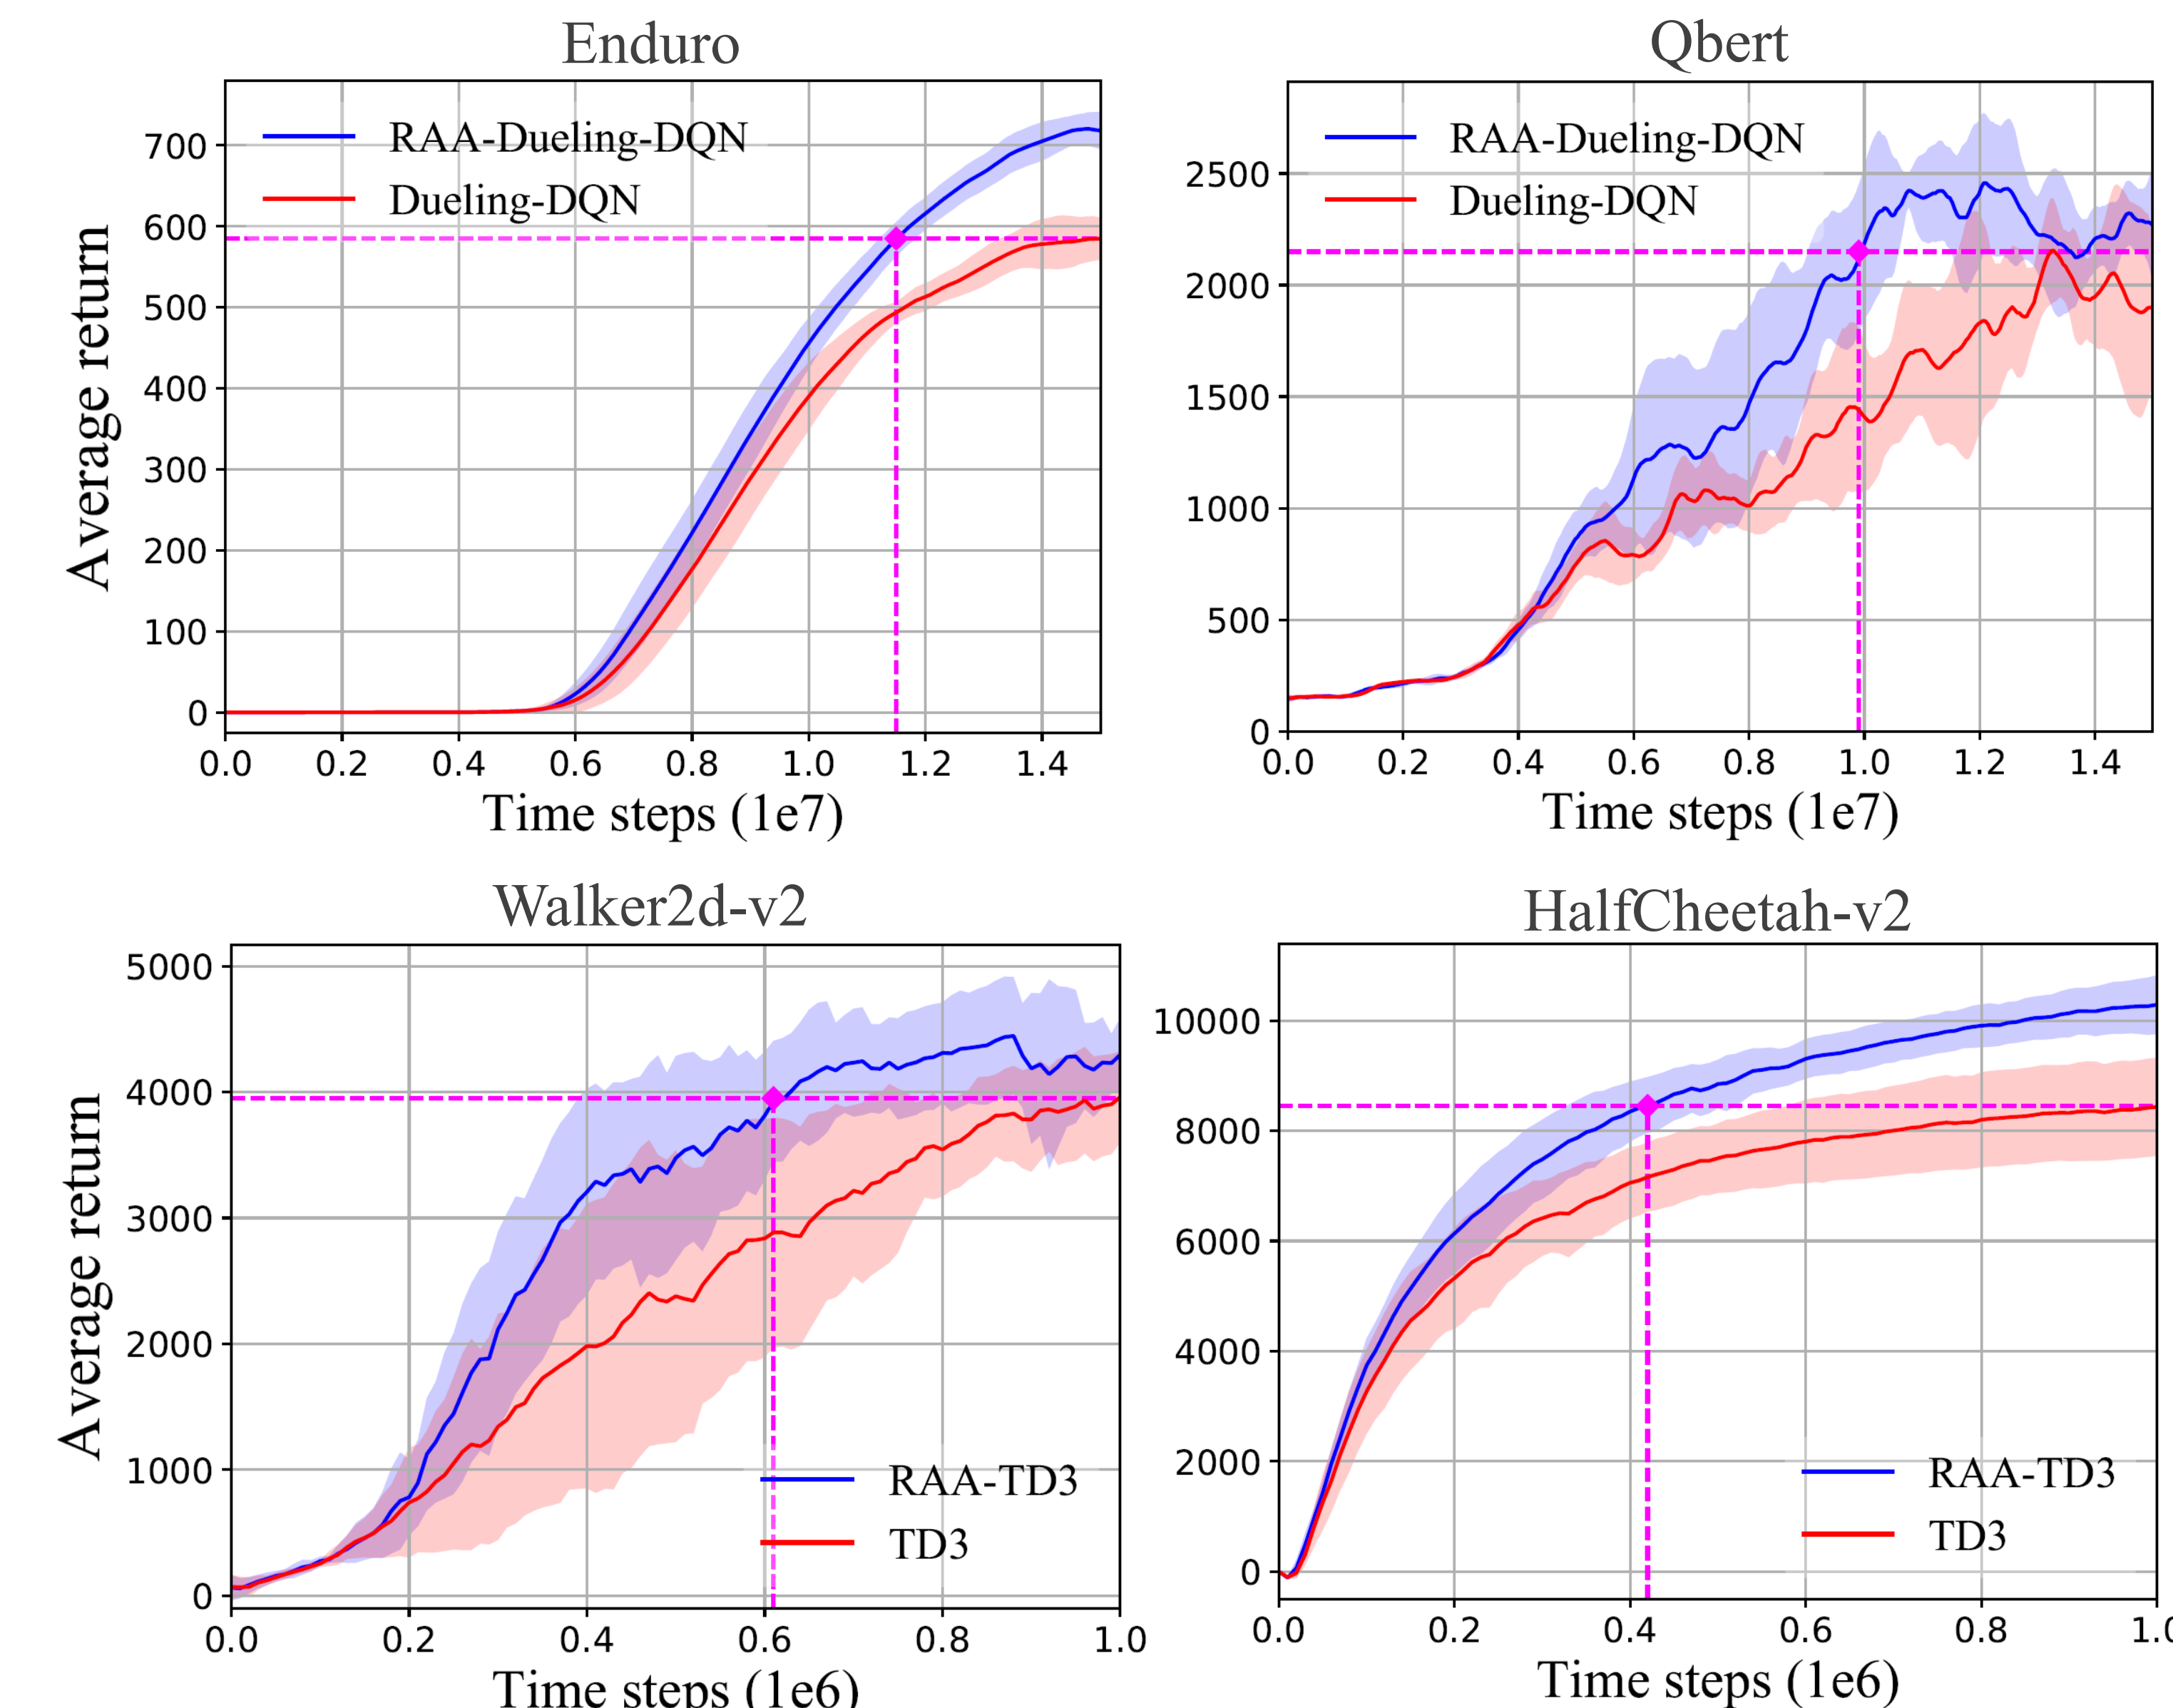


Figure 1: Learning Curves of Dueling-DQN, TD3 and their RAA variants on discrete and continuous control tasks.

Ablation studies

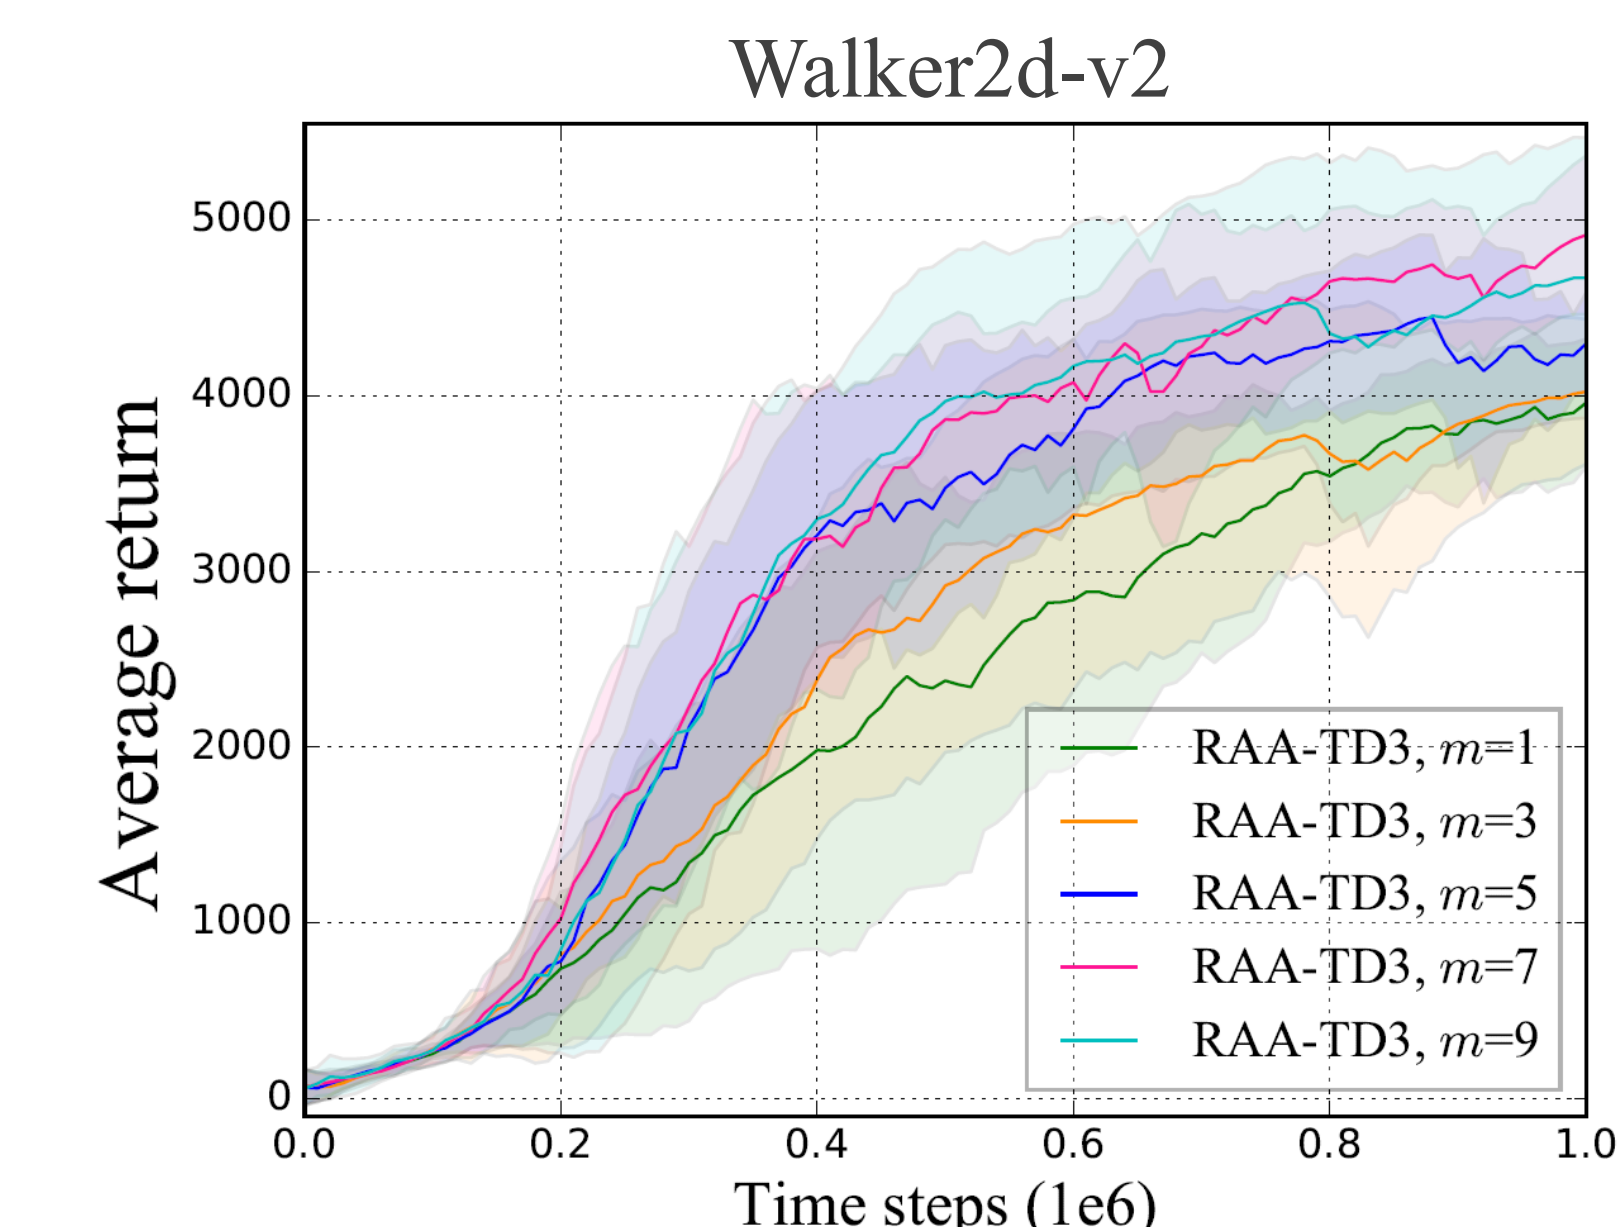


Figure 2: Larger order m leads to faster convergence and better final performance.



GitHub



arXiv